

You are an expert PyTorch researcher implementing a MICCAI-style anatomy-aware vision–language model for radiology.

Your task is to implement the exact architecture described below, without removing or simplifying any component.

📌 Project Overview

Implement an Anatomy-Aware Vision–Language Model trained on the ROCOv2 dataset for both image–text retrieval and image captioning.

The core novelty is:

Structured radiology text encoding (anatomy / observation / modifier)

Anatomy-guided gating of image tokens

Multi-task learning (retrieval + captioning + optional concept prediction)

📁 Codebase Structure (MUST FOLLOW)

Copy code

roco_anatomy_vlm/

```
  └── models/
        ├── image_encoder.py
        ├── text_encoder.py
        ├── anatomy_gating.py
        ├── fusion_transformer.py
        ├── retrieval_head.py
        ├── caption_decoder.py
        └── vlm_model.py

  └── data/
        ├── roco_dataset.py
        └── anatomy_parser.py

  └── training/
        ├── losses.py
        └── train.py

  └── configs/
        └── default.yaml

  └── utils/
        ├── metrics.py
        └── visualization.py
```

① Image Encoder (image_encoder.py)

Use ViT-B/16 from timm

Output patch embeddings only (no classification token)

Shape: (B, N, D)

② Anatomy Parser (anatomy_parser.py)

Rule-based

Input: caption string

Output:

Copy code

```
Python
{
    "anatomy": List[str],
    "observation": List[str],
    "modifier": List[str]
}
```

Use a simple dictionary (RadLex-style keywords)

Must be deterministic and dataset-independent

③ Structured Text Encoder (text_encoder.py)

Use shared BERT weights (bert-base-uncased)

Encode anatomy, observation, modifier separately

Mean-pool each group:

Copy code

Python

```
a_bar, o_bar, m_bar
```

Final text embedding:

Copy code

Python

```
T = concat([a_bar, o_bar, m_bar])
```

④ Anatomy-Guided Gating (anatomy_gating.py)

Implement:

Copy code

Python

```
S_i = cosine_similarity(a_bar, p_i)
```

```
g_i = sigmoid(S_i / temperature)
```

```
p_i_prime = g_i * p_i
```

Must be fully differentiable

Temperature is configurable

⑤ Fusion Transformer (fusion_transformer.py)

Transformer encoder

Input tokens:

Copy code

Python

```
[p_1', ..., p_n', a_bar, o_bar, m_bar]
```

4–6 layers

Output a fused representation

⑥ Retrieval Head (retrieval_head.py)

Projection heads for image and text

CLIP-style contrastive loss

Output Recall@K metrics

⑦ Caption Decoder (caption_decoder.py)

Transformer decoder

Cross-attends to gated image tokens

Teacher forcing during training

⑧ Full Model (vlm_model.py)

Combines all modules

Forward pass supports:

retrieval

captioning

concept prediction (optional)

9 Losses (losses.py)

Implement:

Copy code

Python

$L_{total} = \lambda_1 * L_{retrieval} + \lambda_2 * L_{caption} + \lambda_3 * L_{concept}$

 Implementation Constraints (IMPORTANT)

PyTorch only

No Lightning

Modular, readable, research-quality code

Include docstrings and comments

No shortcuts or architectural simplifications

Assume ROCOv2 captions and images are preprocessed

 Output Expectations

The code must run end-to-end

All tensors must have explicit shapes

All modules must be unit-testable

The architecture must exactly match the description

 Reminder

This is not a generic BLIP or CLIP implementation.

The anatomy-guided gating and structured text encoding are mandatory and central.